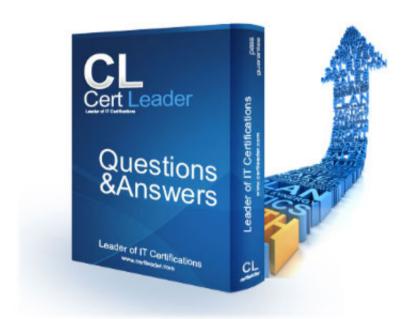# DP-900 Dumps

Microsoft Azure Data Fundamentals

https://www.certleader.com/DP-900-dumps.html

NEW QUESTION 1
To complement the sentence, select the appropriate option in the answer area.

Answer Area

The massively parallel processing (MPP) engine
of Azure Synapse Analytics:
- distributes processing across compute nodes.
- distributes processing across control nodes.
- redirects client connections across compute nodes.
- redirects client connections across control nodes.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Answer Area

The massively parallel processing (MPP) engine
of Azure Synapse Analytics:
- distributes processing across compute nodes.
- distributes processing across control nodes.
- redirects client connections across compute nodes.
- redirects client connections across control nodes.

NEW QUESTION 2
What are two benefits of platform as a service (PaaS) relational database offerings in Azure, such as Azure SQL Database? Each correct answer presents a complete solution. NOTE: Each correct selection is worth one point.

A. complete control over backup and restore processes
B. access to the latest features
C. in-database machine learning services
D. reduced administrative effort for managing the server infrastructure

Answer: BC

NEW QUESTION 3
You need to query a table named Products in an Azure SQL database.
Which three requirements must be met to query the table from the internet? Each correct answer presents part of the solution. (Choose three.)
NOTE: Each correct selection is worth one point.

A. You must be assigned the Reader role for the resource group that contains the database.
B. You must have SELECT access to the Products table.
C. You must have a user in the database.
D. You must be assigned the Contributor role for the resource group that contains the database.
E. Your IP address must be allowed to connect to the database.

Answer: ACE

NEW QUESTION 4
To complete the sentence, select the appropriate option in the answer area.

Batch workloads ▼
- process data in memory, row-by-row.
- collect and process data at most once a day.
- process data as new data is received in near real-time.
- collect data and then process the data when a condition is met.

A. Mastered
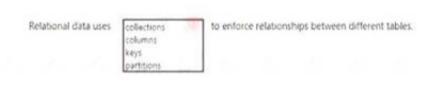B. Not Mastered

Answer: A

Explanation:

Batch workloads

▼

process data in memory, row-by-row.
collect and process data at most once a day.
process data as new data is received in near real-time.
collect data and then process the data when a condition is met.

NEW QUESTION 5
To complete the sentence, select the appropriate option in the answer area.

Answer Area

Relational data uses | collections / columns / keys / partitions | to enforce relationships between different tables.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Answer Area

Relational data uses | collections / columns / keys / partitions | to enforce relationships between different tables.

NEW QUESTION 6
Which Azure SQL offering supports automatic database scaling and automatic pausing of the database during inactive periods?

A. Azure SQL Database Hyperscale
B. Azure SQL managed instance
C. Azure SQL Database serveless
D. Azure SQL Database elastic pod

Answer: B

NEW QUESTION 7
Match the types of activities to the appropriate Azure Data Factory activities.
To answer, drag the appropriate activity type from the column on the left to its Data Factory activity on the right. Each activity type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Activity Types

| Control |
| Data movement |
| Data transformation |

Answer Area

	Copy
	Mapping data flow
	Until

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Activity Types

| Control |
| Data movement |
| Data transformation |

Answer Area

Data movement	Copy
Data transformation	Mapping data flow
Control	Until

NEW QUESTION 8
To complete the sentence, select the appropriate option in the answer area.

You can query a graph database in Azure Cosmos DB	▼
	as a JSON document by using a SQL-like language.
	as a partitioned row store by using Cassandra Query Language (CQL).
	as a partitioned row store by using Language-Integrated Query (LINQ).
	as nodes and edges by using the Gremlin language.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

You can query a graph database in Azure Cosmos DB	▼
	as a JSON document by using a SQL-like language.
	as a partitioned row store by using Cassandra Query Language (CQL).
	as a partitioned row store by using Language-Integrated Query (LINQ).
	as nodes and edges by using the Gremlin language.

NEW QUESTION 9
Match the Azure Data Lake Storage terms to the appropriate levels in the hierarchy.
To answer, drag the appropriate term from the column on the left to its level on the right. Each term may be used once, more than once, or not at all. NOTE: Each correct match is worth one point.

Azure Storage account	File share
File system	

Level

Level

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Terms

Azure Storage account	File share
File system	

Answer Area

Azure Resource Group
└ Azure Storage account
  └ File share
    └Folders
      └Files

NEW QUESTION 10
For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

Statements	Yes	No
Normalization involves eliminating relationships between database tables.	○	○
Normalizing a database reduces data redundancy.	○	○
Normalization improves data integrity.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Statements	Yes	No
Normalization involves eliminating relationships between database tables.	○	⬚○⬚
Normalizing a database reduces data redundancy.	⬚○⬚	○
Normalization improves data integrity.	⬚○⬚	○

NEW QUESTION 10
match the Azure data factory components to the appropriate descriptions.
To answer, drag the appropriate component from the column on the left to its description on the right. Each component may be used once, more than once, or not at all.
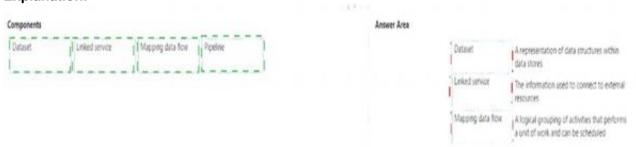NOTE: Each correct match is worth one point.

Components			
Dataset	Linked service	Mapping data flow	Pipeline

Answer Area	
Component	A representation of data structures within data stores
Component	The information used to connect to external resources
Component	A logical grouping of activities that performs a unit of work and can be scheduled

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Components			
Dataset	Linked service	Mapping data flow	Pipeline

Answer Area	
Dataset	A representation of data structures within data stores
Linked service	The information used to connect to external resources
Mapping data flow	A logical grouping of activities that performs a unit of work and can be scheduled

NEW QUESTION 11
For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

Statements	Yes	No
Azure Table storage supports multiple read replicas.	○	○
Azure Table storage supports multiple write regions.	○	○
The Azure Cosmos DB Table API supports multiple read replicas.	○	○
The Azure Cosmos DB Table API supports multiple write regions.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Statements	Yes	No
Azure Table storage supports multiple read replicas.	⬚○⬚	○
Azure Table storage supports multiple write regions.	○	⬚○⬚
The Azure Cosmos DB Table API supports multiple read replicas.	⬚○⬚	○
The Azure Cosmos DB Table API supports multiple write regions.	⬚○⬚	○

NEW QUESTION 15
When provisioning an Azure Cosmos DB account, which feature provides redundancy within an Azure region?

A. multi-master replication
B. Availability Zones
C. automatic failover
D. the strong consistency level

Answer: B

NEW QUESTION 20
Match the security components to the appropriate scenarios.
To answer, drag the appropriate component from the column on the left to its scenario on the right. Each component may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

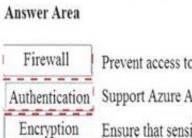| Components	Answer Area	
Authentication		Prevent access to an Azure SQL database from another network.
Firewall		Support Azure Active Directory (Azure AD) sign-ins to an Azure SQL database.
Encryption		Ensure that sensitive data never appears as plain text in an Azure SQL database.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Components	Answer Area	
Authentication	Firewall	Prevent access to an Azure SQL database from another network.
Firewall	Authentication	Support Azure Active Directory (Azure AD) sign-ins to an Azure SQL database.
Encryption	Encryption	Ensure that sensitive data never appears as plain text in an Azure SQL database.

NEW QUESTION 21
You have a transactional application that stem data in an Azure SQ1 managed instance. When should you implement a read-only database replica?

A. You need to generate reports without affecting the transactional workload.
B. You need to audit the transactional application.
C. You need to implement high availability in the event of a regional outage
D. You need to improve the recovery point objective (RPO).

Answer: C

NEW QUESTION 22
When can you use an Azure Resource Manager template?

A. to automate the creation of an interdependent group of Azure resources in a repeatable way
B. to automate the creation of an interdependent and develop can deploy from the Azure portal
C. To provision Azure subscriptions and apply policies for multi-tenant deployments
D. To set complex permissions automatically on existing objects in an Azure subscription

Answer: C

NEW QUESTION 27
For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

Statements	Yes	No
Azure Databricks is an Apache Spark-based collaborative analytics platform.	○	○
Azure Analysis Services is used for transactional workloads.	○	○
Azure Data Factory orchestrates data integration workflows.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Statements	Yes	No
Azure Databricks is an Apache Spark-based collaborative analytics platform.	◎	○
Azure Analysis Services is used for transactional workloads.	○	◎
Azure Data Factory orchestrates data integration workflows.	◎	○

NEW QUESTION 32
To complete the sentence, select the appropriate option in the answer area.

Answer Area

An extract, load, and transform (ELT) process requires

- a data pipeline that includes a transformation engine.
- a separate transformation engine.
- a target data store powerful enough to transform data
- data that is fully processed before being loaded to the target data store.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Answer Area

An extract, load, and transform (ELT) process requires

- a data pipeline that includes a transformation engine.
- a separate transformation engine.
- a target data store powerful enough to transform data
- **data that is fully processed before being loaded to the target data store.**

NEW QUESTION 34
Match the types of data to the appropriate Azure data services.
To answer, drag the appropriate data types from the column on the left to its service on the right. Each data type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Data Types

Image files	Key/value pairs
Relationships between employees	

Answer Area

Data Type	Azure Blob storage
Data Type	Azure Cosmos DB Gremlin API
Data Type	Azure Table storage

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Data Types

Image files	Key/value pairs
Relationships between employees	

Answer Area

Image files	Azure Blob storage
Key/value pairs	Azure Cosmos DB Gremlin API
Relationships between employees	Azure Table storage

NEW QUESTION 37
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

Statements	Yes	No
You can use Azure Data Studio to query a Microsoft SQL Server big data cluster.	○	○
You can use Microsoft SQL Server Management Studio (SSMS) to query an Azure Synapse Analytics data warehouse.	○	○
You can use MySQL Workbench to query Azure Database for MariaDB databases.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Statements	Yes	No
You can use Azure Data Studio to query a Microsoft SQL Server big data cluster.	⊡	○
You can use Microsoft SQL Server Management Studio (SSMS) to query an Azure Synapse Analytics data warehouse.	⊡	○
You can use MySQL Workbench to query Azure Database for MariaDB databases.	⊡	○

NEW QUESTION 42
For each of the following statements, select Yes if the statement is true. Otherwise select No. NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
A pipeline is a representation of a data structure within Azure Data Factory.	○	○
Azure Data Factory pipelines can execute other pipelines.	○	○
A processing step within an Azure Data Factory pipeline is an activity.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Answer Area

Statements	Yes	No
A pipeline is a representation of a data structure within Azure Data Factory.	○	⊡
Azure Data Factory pipelines can execute other pipelines.	⊡	○
A processing step within an Azure Data Factory pipeline is an activity.	○	⊡

NEW QUESTION 43
You have a SQL pool in Azure Synapse Analytics that is only used actively every night for eight hours. You need to minimize the cost of the SQL pool during idle times. The solution must ensure that the data remains intact.
What should you do on the SQL pool?

A. Scale down the data warehouse units (DWUs).
B. Pause the pool.
C. Create a user-defined restore point.
D. Delete the pool

Answer: B

NEW QUESTION 45
You need to use Transact-SQL to query files in Azure Data Lake Storage from an Azure Synapse Analytics data warehouse.
What should you use to query the files?

A. Azure Functions
B. Microsoft SQL Server Integration Services (SSIS)
C. PolyBase
D. Azure Data Factory

Answer: C

NEW QUESTION 50
For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

Statements	Yes	No
You must apply patches to Azure SQL databases regularly.	○	○
You need a Microsoft 365 subscription to create an Azure SQL database.	○	○
You can use existing Microsoft SQL Server licenses to reduce the cost of Azure SQL databases.	○	○

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Statements	Yes	No
You must apply patches to Azure SQL databases regularly.	○	⊙
You need a Microsoft 365 subscription to create an Azure SQL database.	○	⊙
You can use existing Microsoft SQL Server licenses to reduce the cost of Azure SQL databases.	⊙	○

NEW QUESTION 51
To complete the sentence, select the appropriate option in the answer area.

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a [▼]
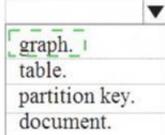
- graph.
- table.
- partition key.
- document.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a [▼]

- graph.
- table.
- partition key.
- document.

NEW QUESTION 54
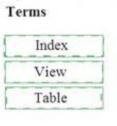Match the terms to the appropriate descriptions.
To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once, or not at all.
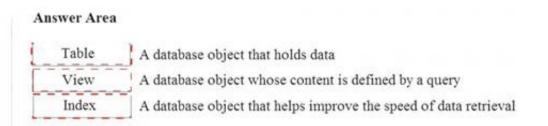NOTE: Each correct match is worth one point.

Terms	Answer Area	
Index		A database object that holds data
View		A database object whose content is defined by a query
Table		A database object that helps improve the speed of data retrieval

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Terms	Answer Area	
Index	Table	A database object that holds data
View	View	A database object whose content is defined by a query
Table	Index	A database object that helps improve the speed of data retrieval

NEW QUESTION 56
Your company plans to load data from a customer relationship management (CRM) system to a data warehouse by using an extract load, and transform (ELT) process.
Where does data processing occur for each stage of the ELT process? To answer, drag the appropriate locations to the correct. Each location may be used once, or not at all, You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.

Locations
An in-memory data integration tool | The CRM system
The data warehouse

Answer Area
Extract: Location
Load: Location
Transform: Location

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Locations
An in-memory data integration tool | The CRM system
The data warehouse

Answer Area
Extract | The data warehouse
Load | An in-memory data integration tool
Transform | The CRM system

NEW QUESTION 58
Your company needs to design a database that shows how changes traffic in one area of a network affect other components on the network.
Which type of data store should you use?

A. Key/value
B. Graph
C. Document
D. columnar

Answer: B

NEW QUESTION 63
At which level in Azure Cosmos DB can you configure multiple write regions and read regions?

A. database
B. partition

C. collection
D. account

Answer: C

NEW QUESTION 67
You need to recommend a non-relational data store that is optimized for storing and retrieving files, videos, audio stream, and virtual disk images. The data store must store data, some metadata, and a unique ID for each file.
Which type of data store should you recommend?

A. document
B. key/value
C. object
D. columnar

Answer: C

NEW QUESTION 71
Your company is designing a data store tor internet-connected temperature sensors. The collected data will be used to analyze temperature trends. Which type of data store should you use?

A. relational
B. columnar
C. graph
D. time series

Answer: A

NEW QUESTION 72
For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

NEW QUESTION 74
Which three objects can be added to a Microsoft Power BI dashboard? Each correct answer presents a complete solution. (Choose three.)
NOTE: Each correct selection is worth one point.

A. a report page
B. a Microsoft PowerPoint slide
C. a visualization from a report
D. a dataflow
E. a text box

Answer: ACE

NEW QUESTION 77
Your company is designing a data store that will contain student data. The data has the following format.

StudentNumber	StudentInformation
7634634	First name: Ben Last: Smith Preferred Name: Benjamin
7634634	First Name: Dominik Last Name: Paiha Email Address: dpaiha@contoso.com MCP ID: 931817
7634636	First Name: Reshma Last Name: Patel Phone number: 514-555-1101
7634637	First Name: Yun-Feng Last Name: Peng

Which type of data store should you use?

A. graph
B. key/value
C. object
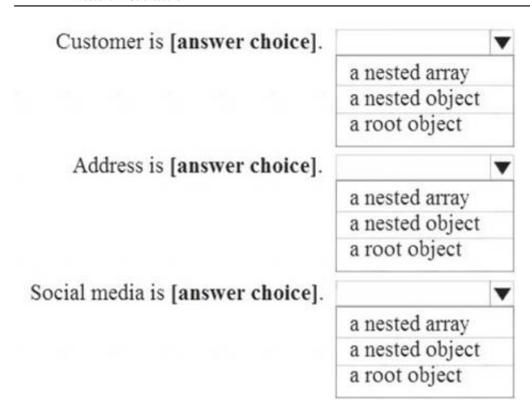D. columnar

Answer: D

NEW QUESTION 79
You have the following JSON document.

```
"customer" : {
    "first name" : "Ben",
    "last name" : "Smith",
    "address" : {
        "line 1" : "161 Azure Ln",
        "line 2" : "Palo Alto",
        "ZIP code" : "54762"
    },
    "social media": [
        {
            "service" : "twitter",
            "handle" : "@bensmith"
        },
        {
            "service" : "linkedin",
            "handle" : "bensmith"
        }
    ],
    "phone numbers": [
        {
            "type" : "mobile",
            "number" : "555-555-555"
        }
    ]
}
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the JSON document.
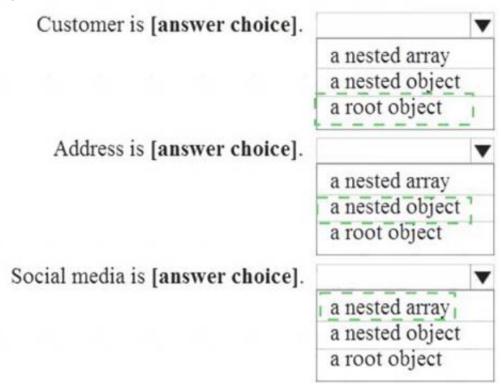NOTE: Each correct selection is worth one point.

Customer is [answer choice].
| a nested array |
| a nested object |
| a root object |

Address is [answer choice].
| a nested array |
| a nested object |
| a root object |

Social media is [answer choice].
| a nested array |
| a nested object |
| a root object |

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Customer is [answer choice].
| a nested array |
| a nested object |
| a root object |

Address is [answer choice].
| a nested array |
| a nested object |
| a root object |

Social media is [answer choice].
| a nested array |
| a nested object |
| a root object |

NEW QUESTION 81
To complete the sentence, select the appropriate option in the answer area.

A relational database must be used when
| a dynamic schema is required. |
| data will be stored as key/value pairs. |
| storing large images and videos. |
| strong consistency guarantees are required. |

A. Mastered
B. Not Mastered

Answer: A

Explanation:

A relational database must be used when
| a dynamic schema is required. |
| data will be stored as key/value pairs. |
| storing large images and videos. |
| strong consistency guarantees are required. |

NEW QUESTION 84
To complete the sentence, select the appropriate option in the answer area.

Answer Area

Relational data is stored in

a file system as unstructured data.
a hierarchial folder structure.
a tabular form of rows and columns.
comma-separated value (CSV) files.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Answer Area

Relational data is stored in

a file system as unstructured data.
a hierarchial folder structure.
a tabular form of rows and columns.
comma-separated value (CSV) files.

NEW QUESTION 89
To complete the sentence, select the appropriate option in the answer area.

In batch processing,

▼
data is always inserted one row at a time.
data is processed in real-time.
latency is expected.
processing car only execute serially.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

In batch processing,

▼
data is always inserted one row at a time.
data is processed in real-time.
latency is expected.
processing car only execute serially.

NEW QUESTION 93
To complete the sentence, select the appropriate option in the answer area.

Answer Area

A Microsoft Power BI

dashboard
interactive report
paginated report
subscription

enables users to create highly formatted,
fixed-layout documents o   rchiving.

A. Mastered
B. Not Mastered

Answer: A

Explanation:
Paginated Reports
https://powerbi.microsoft.com/en-us/blog/announcing-paginated-reports-in-power-bi-general-availability/

NEW QUESTION 97
To complete the sentence, select the appropriate option in the answer area.

Answer Area

A relational database is appropriate for scenarios that involve a high volume of

changes to relationships between entities.
geographically distributed writes.
transactional writes.
writes that have varying data structures.

A. Mastered
B. Not Mastered

Answer: A

Explanation:

Answer Area

A relational database is appropriate for scenarios that involve a high volume of

changes to relationships between entities.
geographically distributed writes.
transactional writes.
writes that have varying data structures.

NEW QUESTION 99
......

Thank You for Trying Our Product

* 100% Pass or Money Back

All our products come with a 90-day Money Back Guarantee.

* One year free update

You can enjoy free update one year. 24x7 online support.

* Trusted by Millions

We currently serve more than 30,000,000 customers.

* Shop Securely

All transactions are protected by VeriSign!

100% Pass Your DP-900 Exam with Our Prep Materials Via below:

https://www.certleader.com/DP-900-dumps.html